

Simulation and Scientific Computing

(Simulation und Wissenschaftliches Rechnen - SiWiR)

Winter Term 2014/15

Florian Schornbaum
Chair for System Simulation



FRIEDRICH-ALEXANDER
UNIVERSITÄT
ERLANGEN-NÜRNBERG

TECHNISCHE FAKULTÄT



Assignment 2: OpenMP-Parallel Red-Black Gauss-Seidel Method

November 3, 2014 – November 25, 2014



FRIEDRICH-ALEXANDER
UNIVERSITÄT
ERLANGEN-NÜRNBERG

TECHNISCHE FAKULTÄT



The Red-Black Gauss-Seidel Method



FRIEDRICH-ALEXANDER
UNIVERSITÄT
ERLANGEN-NÜRNBERG

TECHNISCHE FAKULTÄT

- Elliptic partial differential equation:

$$-\Delta u(x, y) + k^2 u(x, y) = f(x, y)$$

$$-\left(\frac{\partial^2 u(x, y)}{\partial x^2} + \frac{\partial^2 u(x, y)}{\partial y^2}\right) + k^2 u(x, y) = f(x, y)$$


- Discretization using the **differential quotient** for $\Delta u(x, y)$:

$$-\left(\frac{u(x - h_x, y) - 2u(x, y) + u(x + h_x, y)}{h_x^2} + \frac{u(x, y - h_y) - 2u(x, y) + u(x, y + h_y)}{h_y^2}\right) + k^2 u(x, y) = f(x, y)$$

- Discretization using the differential quotient for $\Delta u(x, y)$:

$$-\frac{1}{h_x^2} [u(x - h_x, y) + u(x + h_x, y)] - \frac{1}{h_y^2} [u(x, y - h_y) + u(x, y + h_y)] + \left(\frac{2}{h_x^2} + \frac{2}{h_y^2} + k^2 \right) u(x, y) = f(x, y)$$

- We are solving the given PDE on a **discretized domain** Ω :



$u(x, y) / f(x, y)$

$u(x - h_x, y)$


$u(x + h_x, y)$

$u(x, y - h_y)$

$u(x, y + h_y)$

- For every point $u_{x,y}$ on the grid, we can formulate a linear equation:

$$-\frac{1}{h_x^2} [u_{x-1,y} + u_{x+1,y}] - \frac{1}{h_y^2} [u_{x,y-1} + u_{x,y+1}] + \left(\frac{2}{h_x^2} + \frac{2}{h_y^2} + k^2 \right) u_{x,y} = f_{x,y}$$



$u_{x,y} / f_{x,y}$

$u_{x-1,y}$

$u_{x+1,y}$

$u_{x,y-1}$

$u_{x,y+1}$

- Formulating this equation for every point of the grid leads to a linear system of equations (LSE): $A\vec{u} = \vec{f}$

- The Jacobi method **iteratively** solves the LSE (k represents the number of the iteration) according to the following formula:

$$u_i^{k+1} = \frac{1}{a_{ii}} \left(f_i - \sum_{j \neq i} a_{ij} u_j^k \right)$$

⇒ This corresponds to solving the i -th equation of the LSE using the unknowns from the **previous** iteration.

- Since there are **no data dependencies**, every u_i^{k+1} can be computed in parallel:

```
#pragma omp parallel for
for( int i = 0; i < n; ++i )
    u[i] = ( f[i] - a_west * u_previous[i-1] - a_east * u_previous[i+1]
            - a_south * u_previous[i-m] - a_north * u_previous[i+m]
            ) / a_center;
```

- Solving the i -th equation using the Gauss-Seidel method:

$$u_i^{k+1} = \frac{1}{a_{ii}} \left(f_i - \underbrace{\sum_{j < i} a_{ij} u_j^{k+1}}_{\text{same iteration}} - \underbrace{\sum_{j > i} a_{ij} u_j^k}_{\text{previous iteration}} \right)$$

⇒ Some unknowns come from the previous iteration, some have just been computed in the same iteration $k + 1$.

- Generally **better convergence** compared to the Jacobi method
- There are definitely **data dependencies**: Updating u_i requires some other u_j to be already computed.


- In order to solve the elliptic partial differential equation ...

$$-\Delta u(x, y) + k^2 u(x, y) = f(x, y)$$

... no Matrix A must be assembled.

- We just need to know the “stencil”:

$$\begin{bmatrix} & -\frac{1}{h_y^2} & \\ -\frac{1}{h_x^2} & \frac{2}{h_x^2} + \frac{2}{h_y^2} + k^2 & -\frac{1}{h_x^2} \\ & \frac{1}{h_y^2} & \end{bmatrix}$$



A solution of the PDE

- Every grid point $u_{x,y}$ can be computed as follows:

$$u_{x,y} = \frac{1}{\left(\frac{2}{h_x^2} + \frac{2}{h_y^2} + k^2\right)} \left(f_{x,y} + \frac{1}{h_x^2} [u_{x-1,y} + u_{x+1,y}] + \frac{1}{h_y^2} [u_{x,y-1} + u_{x,y+1}] \right)$$

- Every grid point $u_{x,y}$ is computed as follows (over and over again, iteration for iteration):

$$u_{x,y} = \frac{1}{\left(\frac{2}{h_x^2} + \frac{2}{h_y^2} + k^2\right)} \left(f_{x,y} + \frac{1}{h_x^2} [u_{x-1,y} + u_{x+1,y}] + \frac{1}{h_y^2} [u_{x,y-1} + u_{x,y+1}] \right)$$


⇒ stencil \approx update rule

- Difference between Jacobi and Gauss-Seidel:
 - Jacobi: uses **two grids**, one for storing the values of the previous iteration, one for storing the values of the current iteration
 - Gauss-Seidel: uses only **one grid**, calculates new values “**in place**” (thereby reading some old and some already updated, new values)

- Data dependencies (take a look at the “update rule”/stencil):

$$u_{x,y} = \frac{1}{\left(\frac{2}{h_x^2} + \frac{2}{h_y^2} + k^2\right)} \left(f_{x,y} + \frac{1}{h_x^2} [u_{x-1,y} + u_{x+1,y}] + \frac{1}{h_y^2} [u_{x,y-1} + u_{x,y+1}] \right)$$

$$\begin{bmatrix} & & -\frac{1}{h_y^2} & \\ -\frac{1}{h_x^2} & \frac{2}{h_x^2} + \frac{2}{h_y^2} + k^2 & & -\frac{1}{h_x^2} \\ & & \frac{1}{h_y^2} & \end{bmatrix}$$




North (N)
↕
South (S)

West (W) ↔ East (E)

no data dependencies in NW, NE, SW, and SE direction
data dependencies only in W, E, N, and S direction

- Parallelization → **checkerboard reordering**:




$$\begin{bmatrix} & \text{NORTH} & \\ \text{WEST} & \text{CENTER} & \text{EAST} \\ & \text{SOUTH} & \end{bmatrix}$$

A general stencil for the PDE


- ⇒ Idea: Partition the unknowns into two groups (red and black) so that there are no data dependencies within each group.
- ⇒ Consequence: The unknowns within each group can be updated independently, i.e. in parallel!

- If the red unknowns are updated, only the data of the right-hand side (RHS, $f(x, y)$) that corresponds to those red points is read.
- Splitting the data of the RHS into two separate chunks of memory **improves spatial locality!**
- While the red unknowns are updated, only black unknowns are read, and the red unknowns are only written to memory.
- Many processors have instructions to directly write to memory without first loading the data into the cache: **non-temporal writes**.
- Separating red and black unknowns (just like splitting the data of the RHS) allows for non-temporal writes.
- When splitting red and black data, don't mix up the indices!

- Differences between **SMP** and **NUMA**:



conventional symmetric
multiprocessing (SMP)



non-uniform memory
access (NUMA)

- Memory pages are assigned to the processor that **first touches** the memory – which can be a problem on NUMA systems!
- To avoid non-local memory access, initialize the data by the thread that will later access it and prevent migration of threads to different processors by explicitly **binding/pinning** them.

- Additional material on finite differences by Pascal Frey:

<http://www.ann.jussieu.fr/~frey/cours/UPMC/finite-differences.pdf>

⇒ get it and read it!

- Another well written article on solving linear systems by the same author which might also be worth reading:

<http://www.ann.jussieu.fr/~frey/cours/UPMC/linear%20systems.pdf>



OpenMP



FRIEDRICH-ALEXANDER
UNIVERSITÄT
ERLANGEN-NÜRNBERG

TECHNISCHE FAKULTÄT

- OpenMP ...
 - ... is an **API** that can be used with FORTRAN, C, and C++.
 - ... is a **standard programming model** for shared address space machines.
 - ... allows incremental parallelization.
 - ... provides support for concurrency, synchronization, and data handling.
 - ... obviates the need for explicitly setting up mutexes and condition variables.
- The workload is distributed among threads:
 - Variables can be ...
 - ... shared among all threads or ...
 - ... duplicated for each thread.
 - **Threads communicate by sharing variables.**

- Setting the total number of threads:

- At runtime via the function “omp_set_num_threads”:

```
#include <omp.h>
void omp_set_num_threads( int num_threads )
```

- Via the environment variable OMP_NUM_THREADS

```
export OMP_NUM_THREADS=4 (if you are using bash)
```

- Getting the total number of threads:

```
#include <omp.h>
int omp_get_num_threads( void )
```

- Getting the ID of the thread that executes this command:

```
#include <omp.h>
int omp_get_thread_num( void )
```

- OpenMP directives in C and C++ are based on **#pragma** compiler directives:

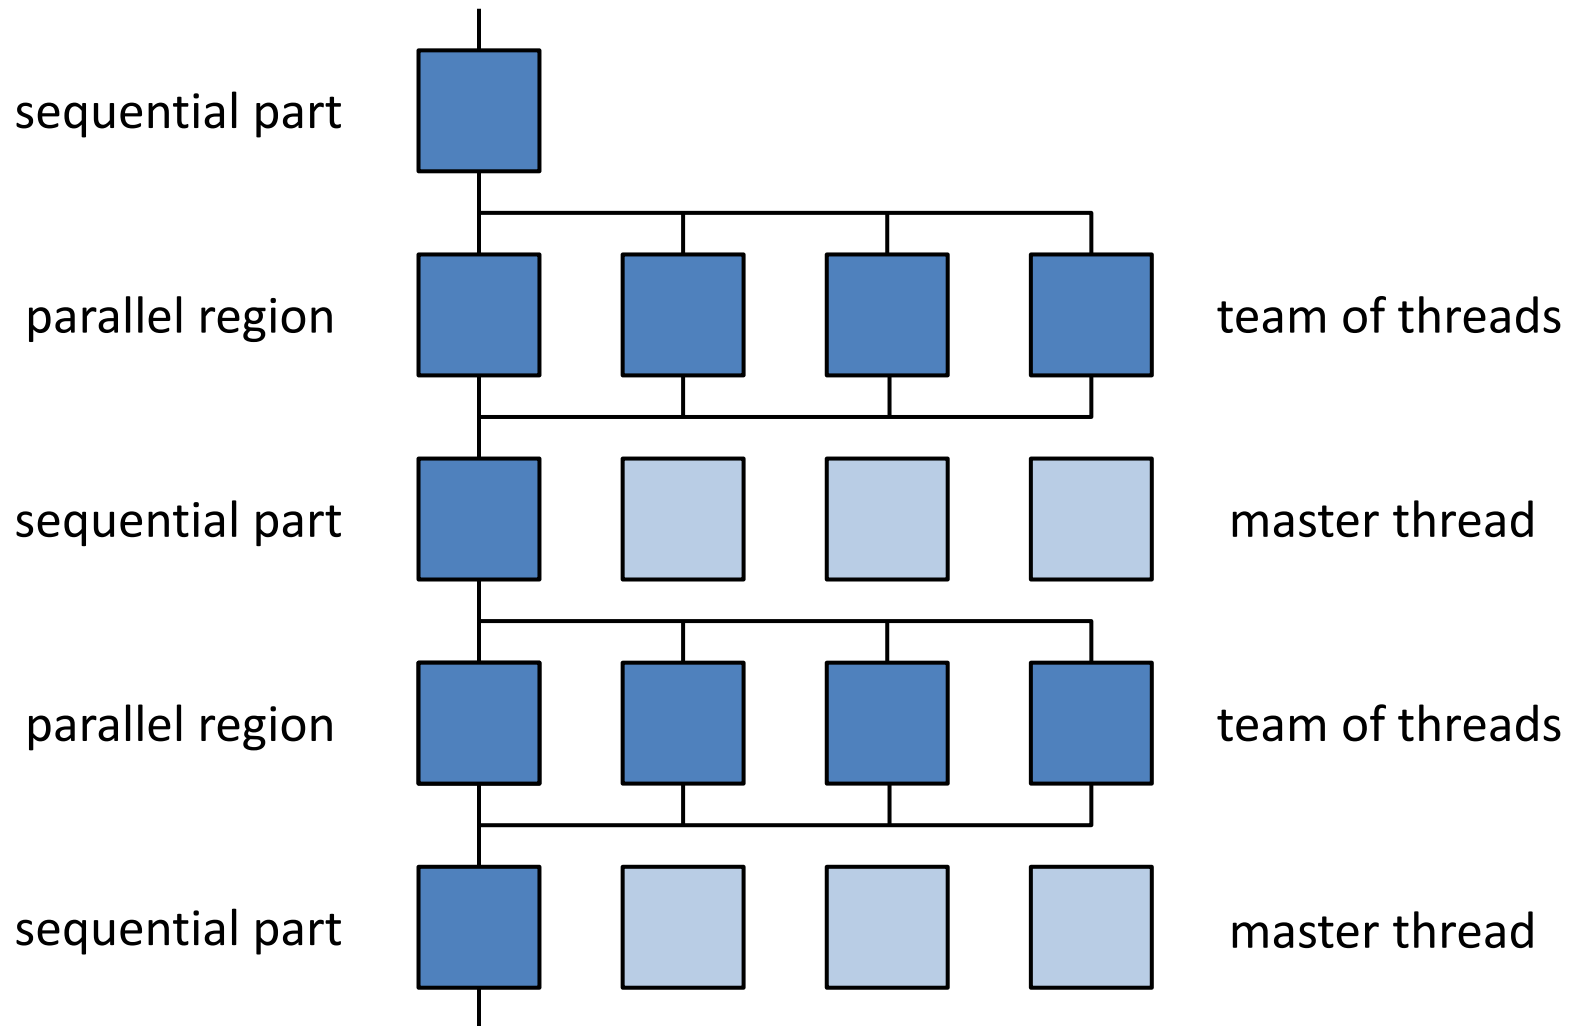
```
#pragma omp directive [clause list]
```

- OpenMP programs execute serially until they encounter the **parallel** directive:

```
#pragma omp parallel [clause list]  
/* structured block */
```

- Each thread executes the structured block specified by the **parallel** directive.
- The **clause list** specifies the conditional parallelization, the number of threads, and data handling.

The OpenMP Execution Model



```
#include <omp.h>
#include <iostream>

int main ()
{
    /* Fork a team of threads */
    #pragma omp parallel
    {
        /* Obtain and print the thread id */
        int tid = omp_get_thread_num();
        std::cout << "Hello world from thread " << tid << std::endl;

        /* Only the master thread does this */
        if( tid == 0 ) {
            int nthreads = omp_get_num_threads();
            std::cout << "Number of threads = " << nthreads << std::endl;
        }

    } /* All threads join the master thread and terminate */

    return 0;
}
```

- The clause `if(scalar expression)` determines whether the parallel construct results in the creation of threads (only one if clause can be used with a parallel directive).

```
#pragma omp parallel if( is_parallel == 1 )  
/* structured block */
```

- The clause `num_threads(integer expression)` specifies the number of threads that are created by the parallel directive.

```
#pragma omp parallel num_threads( 8 )  
/* structured block */
```

- **private(variable list)**: indicates that the set of variables is local to each thread – i.e., each thread gets its own copy of each variable in the list. Variables declared **private** should be assumed to be uninitialized for each thread.
- **firstprivate(variable list)**: same as **private**, however, all variables are initialized with copies of the values of the original variables.
- **shared(variable list)**: indicates that all variables in the list are shared across all the threads.
 - Shared variables exist in only one memory location – all threads read and write to that address (that's how threads can communicate).
 - It's your responsibility that multiple threads properly access these variables.

```
// The parallel section is executed by 8 threads
// depending on the value of 'is_parallel'. Each of
// the threads gets a local copy of a and c, and
// shares the value of b. Furthermore, the value of
// each copy of c is initialized to the value of c
// before the parallel directive.

#pragma omp parallel if( is_parallel == 1 ) \
                    num_threads( 8 ) private( a ) \
                    shared( b ) firstprivate( c )
{
    // parallel section
}
```


- **reduction**: specifies how multiple “local copies” of a shared variable are combined into a single copy for the master thread when the parallel section exits.
 - possible operators: +, -, *, &, |, ^, &&, and ||.

```
double sum = 0.0;
#pragma omp parallel shared( sum ) reduction(+: sum) \
                    num_threads( 8 )
{
    // parallel section, executed by 8 threads
    /* thread local sum is computed here */
}
/* here, sum contains the sum of all local instances of sum */
```

- Work-sharing constructs ...
 - ... must be enclosed within a parallel region.
 - ... divide the execution of the enclosed code region among all currently existing threads.
 - ... do not launch new threads.
 - ... have no implicit barrier on entry.
- Work-sharing constructs are:
 - The **for** directive
 - The **sections** directive

- The `for` directive specifies that the iterations of the loop immediately following it must be executed in parallel by all currently existing threads.
- The general form looks like as follows:

```
#pragma omp for [clause list]
/* for loop */
```
- The `for` directive can be fused with the `parallel` directive:

```
#pragma omp parallel for [clause list]
/* for loop */
```
- Possible clauses: `private`, `firstprivate`, `shared`, `reduction`
(The loop control variable is private by default)

The for Directive - Example


```
int a[20];  
int b[20];  
int f = 23;  
  
/* b is initialized */  
  
#pragma omp parallel for schedule( static )  
for( int i = 0; i < 20; ++i )  
    a[i] = b[i] + f*(i+1)
```

If this program is compiled, `OMP_NUM_THREADS` is set to 4, and then the program is executed, each thread calculates exactly 5 values in the array `a`. The first thread computes `a[0]` to `a[4]`, the second thread `a[5]` to `a[9]`, ...

- `schedule(scheduling_class[, optional parameters])`:
 - `static[, chunk_size]`: Loop iterations are divided into pieces of size `chunk` and then assigned to all threads in a round robin fashion. If no chunk size is specified, the iterations are evenly divided among all threads.
 - `dynamic[, chunk_size]`: Loop iterations are divided into pieces of size `chunk` and assigned to threads as they become idle (= dynamic scheduling). The default chunk size is 1.
 - `guided[, chunk_size]`: Similar to dynamic except that the block size decreases each time work is given to a thread. Initial block size is equal to $\frac{\#iterations}{\#threads}$, subsequent blocks are proportional to $\frac{\#remaining_iterations}{\#threads}$. The chunk size defines the minimum block size (default is 1).
 - `runtime`: strategy determined by environment variable `OMP_SCHEDULE`
- Next slide: example with 12 iterations and 3 threads

The for Directive - Scheduling

w/o OpenMP



static



dynamic, 3



guided, 1




- There is an implicit barrier at the end of the loop – specify **`nowait`** in order omit this barrier.
- It is illegal to branch/**`break`** out of a parallel for loop.
- The loop control variable must be ...
 - ... a signed integer.
 - ... an unsigned integer (OpenMP 3.0)
 - ... a pointer type (OpenMP 3.0)
 - ... a random access iterator (C++ only, OpenMP 3.0)
- The number of iterations must be computable at loop entry.
- The logical expression must be one of `<`, `<=`, `>`, `>=`.
- The increment expression must have integer increments or decrements only.

- The **sections** directive is used for non-iterative parallel task assignment. In general it looks like as follows:

```
#pragma omp sections [clause list]
{
    #pragma omp section
    {
        /* first section */
    }
    #pragma omp section
    {
        /* second section */
    }
    ...
}
```


The sections Directive – Example

```
#pragma omp parallel {  
  #pragma omp sections  
  {  
    #pragma omp section  
    {  
      a = ...;  
      b = ...; ...; }  
  
    #pragma omp section  
    {  
      c = ...;  
      d = ...; ...; }  
  
    #pragma omp section  
    {  
      e = ...;  
      f = ...; ...; }  
  
    #pragma omp section  
    {  
      g = ...;  
      h = ...; ...; }  
  }  
}
```



The sections Directive – Example

```
#include <omp.h>

int main ()
{
    int a[20], b[20], c[20], d[20];

    /* initialization of array a and b */

    #pragma omp parallel sections
    {
        #pragma omp section
        {
            for( int i = 0; i != 20; ++i )
                c[i] = a[i] + b[i];
        }

        #pragma omp section
        {
            for( int i = 0; i != 20; ++i )
                d[i] = a[i] * b[i];
        }
    }

    return 0;
}
```

The **sections** directive
can be combined with
the **parallel** directive.

- The **barrier** directive synchronizes all currently active threads:

```
#pragma omp barrier
```

- The **critical** directive specifies a region of code that ...
 - ... is executed by all threads.
 - ... must be executed by only one thread at a time.

```
#pragma omp critical [(optional name)]  
/* structured block */
```

A thread waits at the beginning of a critical region until no other thread in the team is executing a critical region with the same name. All unnamed critical directives map to the same unspecified name.

- The **master** directive specifies a region of code that is only executed by the master thread – all other threads skip this region.

```
#pragma omp master  
/* structured block */
```

There are no implicit barriers!

- The **single** directive specifies a region of code that is only executed by one thread (not necessarily the master) – all other threads skip this region.

```
#pragma omp single [clause list]  
/* structured block */
```

There is an implicit barrier at the end of the block – unless a **nowait** clause is specified.

```
int a[20], b[20];
int c = 0, f = 23;

/* b is initialized */

#pragma omp parallel
{
    #pragma omp for schedule( static )
    for( int i = 0; i < 20; ++i )
    {
        if( b[i] == 0 )
        {
            #pragma omp critical
            ++c;
        }
        a[i] = b[i] + f*(i+1)
    }

    #pragma omp single
    std::cout << "zeros in b: " << c << std::endl;
}
```

- A well written and quite comprehensive overview on OpenMP:

<https://computing.llnl.gov/tutorials/openMP/>

⇒ definitely worth reading!

- The official OpenMP API specification:

<http://openmp.org/wp/openmp-specifications/>

THE END QUESTIONS ?



FRIEDRICH-ALEXANDER
UNIVERSITÄT
ERLANGEN-NÜRNBERG

TECHNISCHE FAKULTÄT